

ROSSO HAN

Engineering Manager | Applied AI, Agent Systems & RAG

Oakland, CA | +1 (510) 458-9626 | jian.han3@gmail.com |
linkedin.com/in/rossohan

PROFILE

Engineering manager and hands-on Applied AI systems builder with 15+ years across backend platforms, cloud architecture, APIs, distributed systems, and technical leadership. Builds agentic AI and RAG products with memory, guardrails, evals, observability, human control, and reliable fallbacks. Combines production engineering judgment with founder-level execution from problem discovery through architecture, implementation, testing, and validation.

CORE STRENGTHS

Applied AI: agentic workflows, RAG, LLM orchestration, AI memory, prompt systems, guardrails, evals, semantic search, vector databases

Engineering Leadership: team development, architecture, roadmaps, execution, incident review, cross-functional alignment, distributed teams

Platform Engineering: APIs, microservices, AWS, reliability, CI/CD, Terraform, observability, performance, production operations

APPLIED AI WORK

Founder & Applied AI Systems Builder | TryForgeAI | 2026 - Present | Oakland, CA

- **Forge AI RAG:** Built a local mixed-format knowledge system that ingests Markdown, Excel, and Word files, creates Voyage AI embeddings, persists 487 indexed chunks in Chroma, retrieves relevant evidence, and generates grounded answers through Claude.
- **AI automation:** Built a nine-node n8n workflow connecting scheduled execution, Google Sheets CRM data, lead filtering, Claude API reasoning, CRM updates, and Gmail notifications; created reusable skills for market intelligence and prospecting.

- **Avaloka AI:** Implemented risk routing, LLM orchestration, guardian/repair/fallback, versioned prompts, and a SAGE-style memory writer, guardian, Care Card store, deterministic reader, lifecycle controls, diagnostics, and eval harness. Verified with 77 passing tests.
- **Auralith:** Built a modular AI-to-IoT pipeline that maps emotional context into guarded responses and safe low-stimulation light behavior, with AI shadow mode, scene validation, WLED adaptation, local-network controls, and a design-partner Pilot Kit. Verified with 200 passing tests across six modules.
- **Agent-first delivery:** Created durable project operating systems that make product direction, architecture decisions, task state, traces, evals, and validation recoverable by humans and coding agents.

PROFESSIONAL EXPERIENCE

Engineering Manager | Rakuten Rewards | 2015 - Present

- Led engineering teams responsible for backend and API infrastructure supporting a high-traffic consumer rewards platform.
- Managed distributed collaboration across five time zones and aligned engineering, product, operations, and business stakeholders around delivery and reliability.
- Guided architecture, planning, technical review, CI/CD practices, incident response, and continuous improvement for business-critical services.
- Improved platform performance and resilience through caching, load balancing, microservice optimization, observability, and disciplined production operations.
- Mentored engineers and translated ambiguous business and data-flow problems into executable technical plans.

Cloud Architect | Hewlett-Packard (HP) | 2010 - 2015

- Designed cloud-oriented APIs, service architecture, and resilient infrastructure for enterprise systems.
- Supported modernization from monolithic applications toward scalable, maintainable services and more reliable data flows.

- Introduced Infrastructure as Code practices with Terraform and improved repeatability across environment setup and deployment.
- Worked across engineering and stakeholder groups to improve cloud readiness and operational reliability.

LLM & ENTERPRISE RAG

SupportVectors LLM and Enterprise RAG Bootcamp, Summer 2026. Applied study of embeddings, chunking, dense and sparse retrieval, fusion, reranking, query transformation, grounding, guardrails, RAG evaluation, RAPTOR, GraphRAG, semantic caching, SAGE precise retrieval, and Text-to-SQL. Course concepts are translated into Avaloka experiments, trace schemas, eval cases, and architecture decisions; advanced methods are treated as research until measured against a baseline.

TECHNICAL TOOLBOX

AI / Agents: OpenAI API, Anthropic Claude, LLM orchestration, RAG, Voyage AI, Chroma, semantic search, vector retrieval, memory systems, prompt registries, evals, guardrails

Automation: n8n, agent skills, scheduled workflows, Google Sheets API, Gmail integrations, human-in-the-loop approval

Engineering: Python, TypeScript, Java, React, Node.js, REST APIs, microservices, distributed systems

Cloud / Delivery: AWS, Terraform, CI/CD, observability, incident management, caching, load balancing

EDUCATION & CERTIFICATIONS

M.S., E-Business Management - University of Warwick, UK

B.S., Computer Science - Beijing University of Technology, China

Certifications: PMP; Certified ScrumMaster (CSM)